

Solution to Series 9

1. a) Read in the data:

```
> ## load and transform the data:
> data(ozone, package = "gss")
> d.ozone <- subset(transform(ozone, logupo3 = log(upo3)), select = -upo3)
> d.ozone.e <- d.ozone[~which.max(d.ozone[, "wdsp"]),]
> d.ozone.es <- d.ozone.e
> d.ozone.es[, -10] <- scale(d.ozone.e[, -10])
```

- b) We fit the MARS model with maximal interaction degree 2.

```
> library(earth) ## for MARS
> mars.fit2 <- earth(formula = logupo3 ~ ., data = d.ozone.e, degree = 2 )
> summary(mars.fit2, digits=3)
```

Call: earth(formula=logupo3~., data=d.ozone.e, degree=2)

	coefficients
(Intercept)	2.7570
h(sbtp-52)	0.0197
h(13-dgpg)	-0.0091
h(dgpg-13)	-0.0062
h(194-ibtp)	-0.0053
h(200-vsty)	0.0014
h(66-day)	-0.0130
h(day-66)	-0.0025
h(8-wdsp) * h(200-vsty)	0.0002
h(wdsp-8) * h(200-vsty)	-0.0029
h(hmdt-71) * h(52-sbtp)	-0.0051
h(hmdt-51) * h(ibht-1049)	0.0000
h(1049-ibht) * h(day-286)	0.0000
h(1049-ibht) * h(286-day)	0.0000

Selected 14 of 21 terms, and 8 of 9 predictors

Termination condition: Reached nk 21

Importance: sbtp, ibtp, day, dgpg, vsty, ibht, hmdt, ...

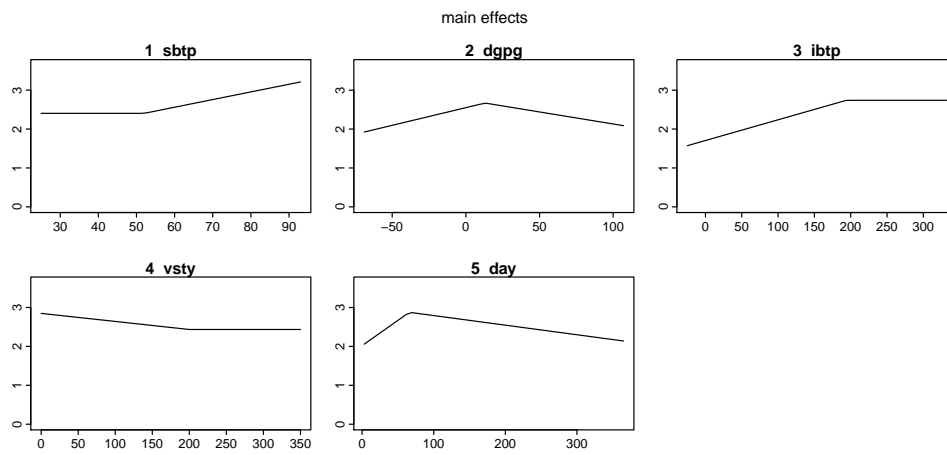
Number of terms at each degree of interaction: 1 7 6

GCV 0.107 RSS 28.5 GRSq 0.809 RSq 0.845

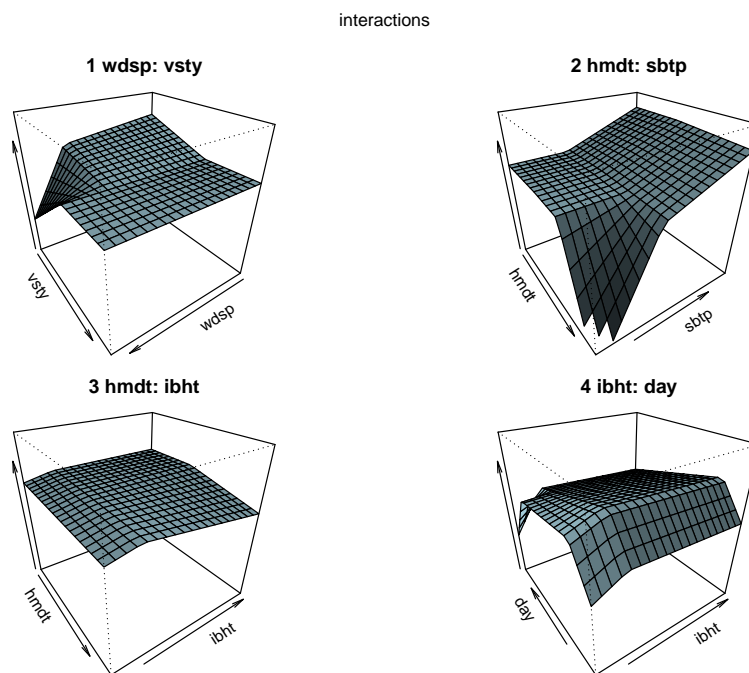
```
> ## Plot the main effects
```

```
> plotmo(mars.fit2, degree2=FALSE, caption="main effects")
```

```
grid:      vdht wdsp hmdt sbtp ibht dgpg ibtp vsty day
          5760   5   64   62 2109   24  168  120 178
```

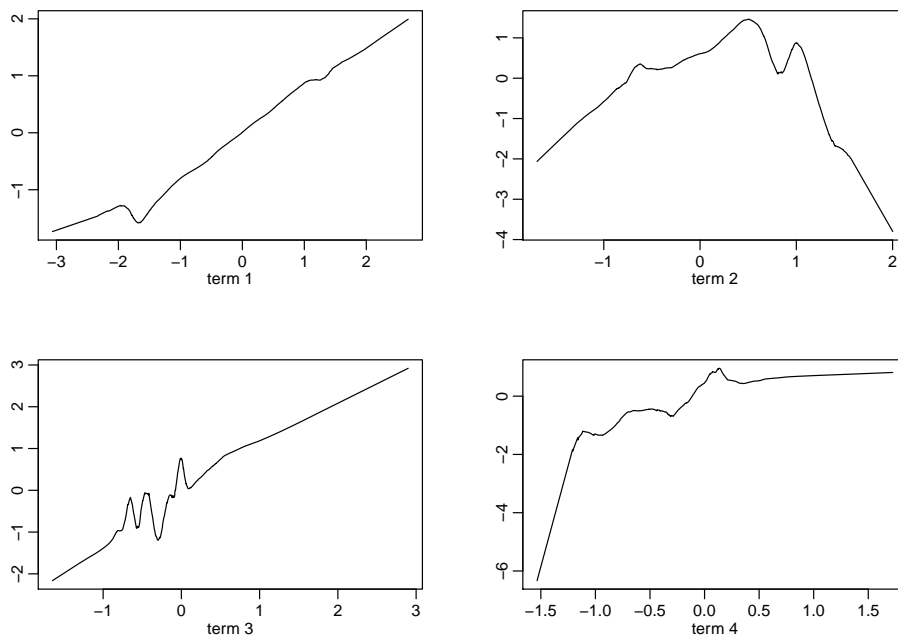


```
> ## Plot the interaction
> plotmo(mars.fit2, degree1=FALSE, caption="interactions")
```



c) Basic PPR fit and plot:

```
> ppr.fit <- ppr(logupo3~, nterms=4, data=d.ozone.es)
> sfsmisc::mult.fig(4)
> plot(ppr.fit, type='l')
```



The fitted lines of the ridge functions are too wiggly, therefore the model is not optimal. A better model would have smoother fitted lines of the ridge functions.

- d) The following R-code calculates the CV-error for $\text{nterms} \in \{3, 4, 5\}$ and all three smoothers $\text{sm.method} \in \{\text{"supsmu"}, \text{"spline"}, \text{"gcv spline"}\}$. $\text{df} \in \{5, 6, 7\}$ are used for spline. max.terms is always set equal to $\text{nterms} + 3$.

```
> ## load the cv function:
> source("http://stat.ethz.ch/education/semesters/ss2014/CompStat/Exercises/cv-fun.R")
> ## number of terms
> N<- 3:5
> ## number of degrees of freedom
> DF<- 5:7
> ## Initialize cv scores
> cv.sm <- numeric(length = length(N))
> cv.gcv <- numeric(length = length(N))
> cv.ss <- matrix(0, nrow = length(N), ncol = length(DF),
+                 dimnames=list(paste('df',DF), paste('nterms',N)))
> for(i in seq_along(N)){
+   n <- N[i]
+   cv.sm[i] <- cv(fitfn=ppr, nterms = n, max.terms = n + 3)
+   cv.gcv[i] <- cv(ppr, sm.method = "gcv spline", nterms = n, max.terms = n + 3)
+   for(j in seq_along(DF)){
+     df <- DF[j]
+     cv.ss[i, j] <- cv(ppr, sm.method = "spline", df = df, nterms = n, max.terms = n + 3)
+   }
+ }
```

- e) Fit MARS models and print their cv:

```
> mars.1 <- cv(earth, degree = 1)
> mars.2 <- cv(earth, degree = 2)
> mars.3 <- cv(earth, degree = 3)
```

Compare with PPR cv:

```
> ##print the cv scores of the PPR fits
> cv.sm
[1] 53.05948 54.07697 53.92102
> cv.gcv
[1] 53.17953 54.75474 58.35856
```

```

> cv.ss
      nterms 3 nterms 4 nterms 5
df 5 40.23206 38.70816 43.42077
df 6 38.76893 35.00961 43.22224
df 7 38.38503 36.25633 44.19653

> ##print the cv scores of the MARS fits
> c(cv.mars.1=mars.1,cv.mars.2=mars.2 ,cv.mars.3=mars.3)

cv.mars.1 cv.mars.2 cv.mars.3
40.06201  43.95574  44.21432

```

The best model in this case uses `sm.method = "spline"`, `nterms = 4` and `df = 6`. Please note that there are other options to “tune” `supsmu` and `gcv spline` which should also be considered to do a fair comparison (see `bass`, `span`, `gcvpen`). MARS with interaction degree of 1,2 and 3 results in `cv-errors` of 41.14, 43.76 and 43.95, respectively. These values are worse than the best solution found with Projection Pursuit Regression (35.01).

Now let's compute the CV-error of a Generalized additive model (GAM). Here we need to write a wrapper for the training algorithm, otherwise we run into trouble with the special formulas of `gam()`.

```

> library(mgcv)
> wrapGam <- function(formula, data,...){
  gamForm <- wrapFormula(formula, data = data)
  gam(gamForm, data=data, ...)
}
> gam.cv <- cv(wrapGam, formula=as.formula('logupo3~.'), trace=FALSE)

```

The CV-error of GAM is 37.98, which is bigger than the best PPR, but compete quite well with all the fitted models.

- f) The 4 ridge functions look much smoother than those of task c). In the former we used the super smoother (default value), which adapts the smoothness internally. The optimal solution found by `cv`, however, uses a spline with $df = 6$. The df of the spline fit are smaller than the “effective df ” of the super smoother fit. Therefore, the spline fit looks smoother.

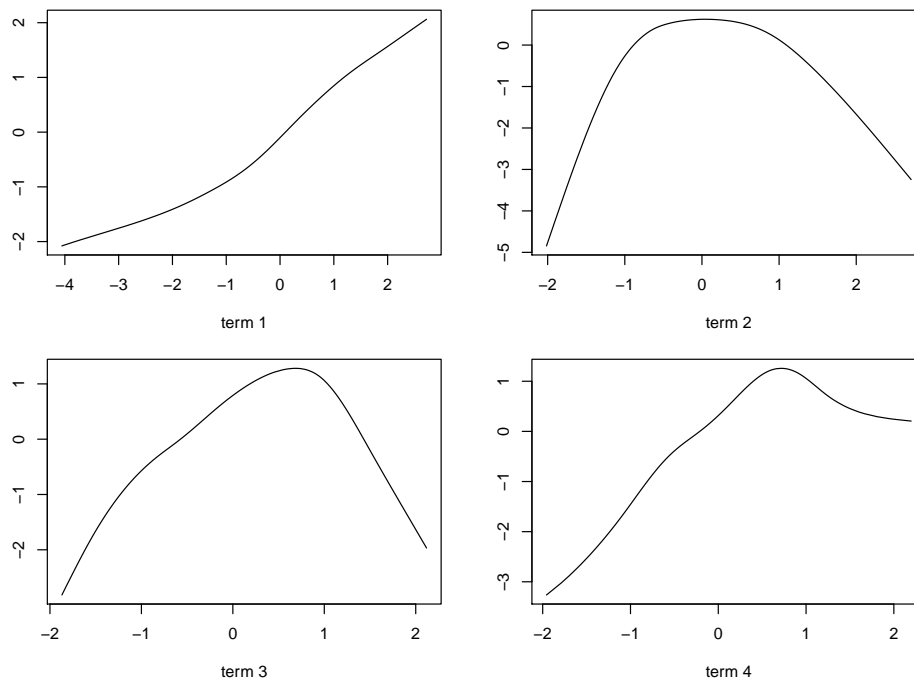
```

> ## Which index corresponds to min(cv.ss)?
> which(cv.ss==min(cv.ss),arr.ind=TRUE)

      row col
df 6    2   2

> fit <- ppr(formula=logupo3~.,data=d.ozone.es, sm.method = "spline",
             df = 6, nterms = 4, max.terms = 7)
> par(mfrow = c(2,2))
> plot(fit, type='l')

```



g) The scaling factors (β_k in the manuscript) of the ridge terms are

```
> round(fit$beta,digits=3)
term 1 term 2 term 3 term 4
0.534  0.263  0.312  0.305
```

and the α -matrix is given by

```
> round(fit$alpha,2)
      term 1 term 2 term 3 term 4
vdht   0.48  -0.09  -0.06  -0.02
wdsp  -0.25  -0.11   0.09   0.12
hmdt  -0.03  -0.54   0.16   0.22
sbtp   0.48   0.43  -0.04  -0.37
ibht  -0.06  -0.17  -0.05   0.01
dgpg   0.47  -0.45  -0.11  -0.66
ibtp   0.07  -0.10  -0.10   0.32
vsty  -0.44  -0.10   0.05   0.26
day   -0.23  -0.49  -0.97   0.44
```

The coefficient of the variable `ibht` has relatively small absolute values in all of the terms: It seems to be not so important. `term 3` is dominated mainly by the value of `day`.